

# Going Beyond the Basics: Advanced Deep Learning Techniques with TensorFlow

**Gusti Triandi Winata**

Community Researcher of RL & LLM  
Cohere for AI

## About Me



# Gusti Triandi Winata

### Latest Work Experiences:

- |                                                   |                |
|---------------------------------------------------|----------------|
| • Researcher, Cohere for AI                       | 2024 - present |
| • MLE Consultant, Freeport Indonesia              | 2023 - 2024    |
| • Mid. MLE, eFishery                              | 2022 - 2023    |
| • OSS Fellowship with Adobe, Major League Hacking | 2021           |

### Education:

- |                                                                                                                                                         |                   |
|---------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|
| • Bandung Institute of Technology<br><i>Bachelor of Electrical and Computer Engineering</i><br><i>Was a Bangkit Graduate of the first batch (2020)!</i> | Graduated at 2021 |
|---------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|

# Ground Rules

Observe the following rules to ensure a supportive, inclusive, and engaging classes



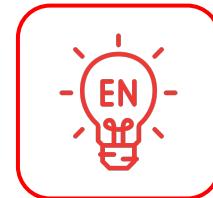
Give full attention  
in class



Mute your microphone  
when you're not talking



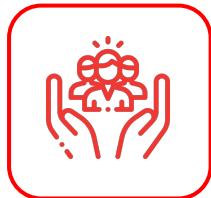
Keep your  
camera on



Turn on the CC Feature  
on Meet

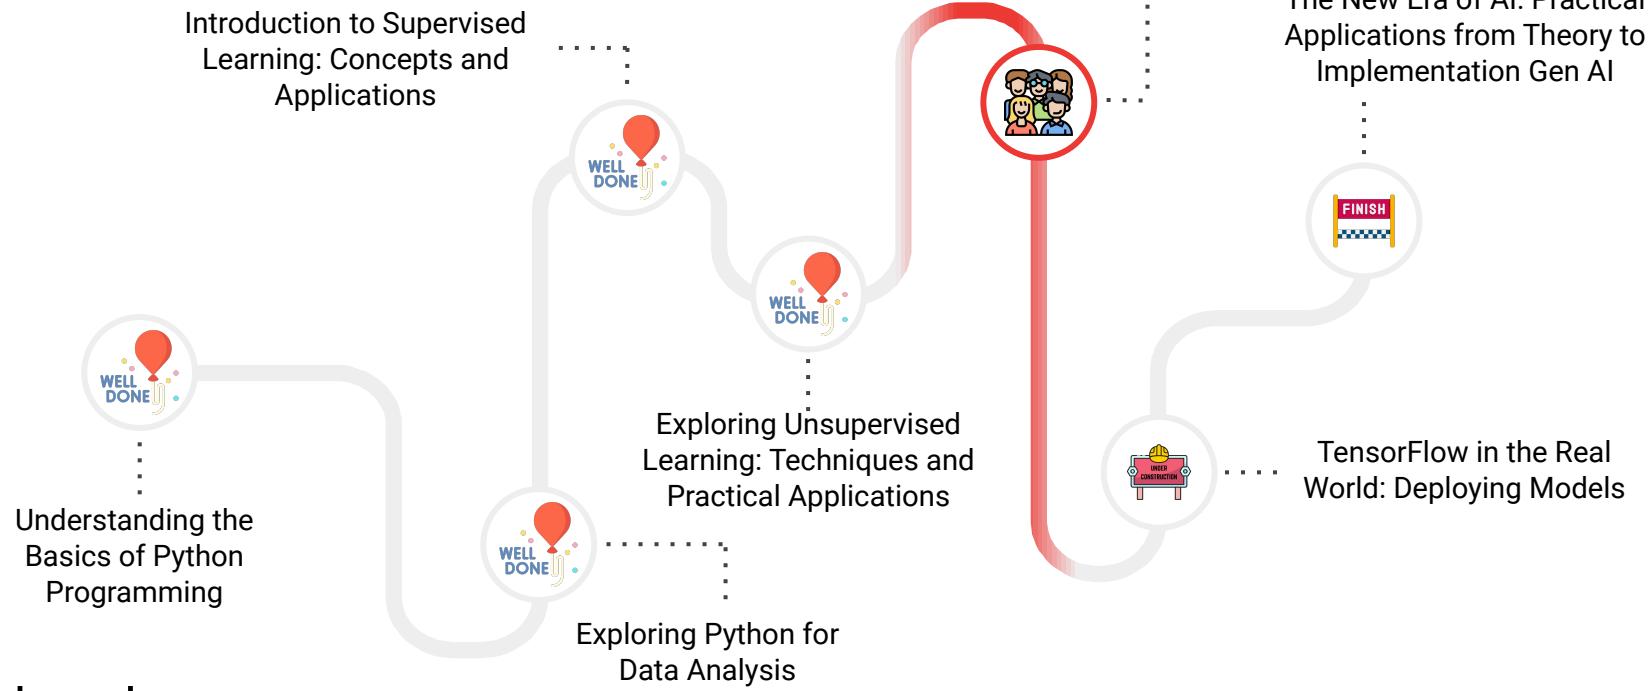


Use raise hand or chat  
to ask questions

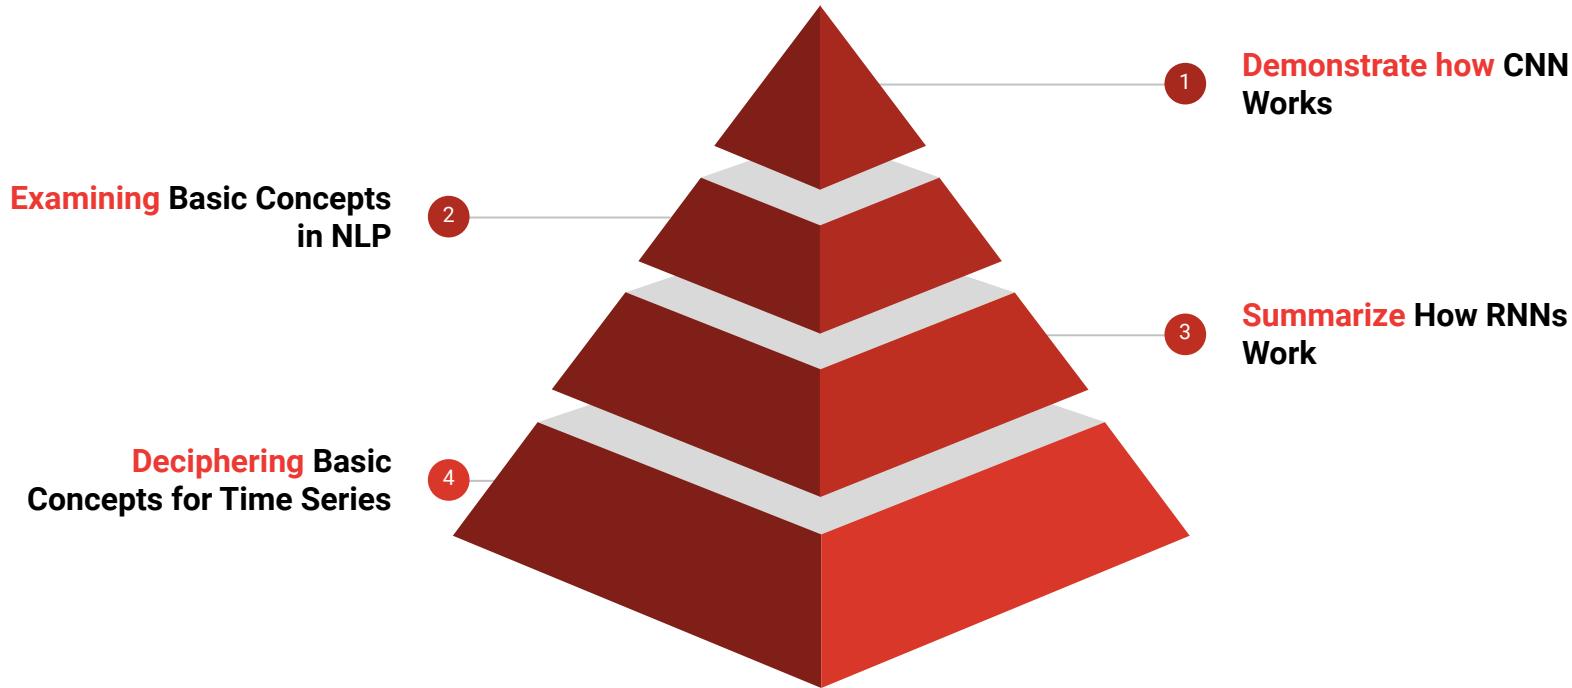


Make this room a safe place  
to learn and share

# Where Are We Now?



# Learning Objectives



# Today's Agenda

**Image Classification with  
Convolutional Neural  
Network**



1



2

**An Overview of Natural  
Language Processing**

**Exploring RNNs and LSTM**



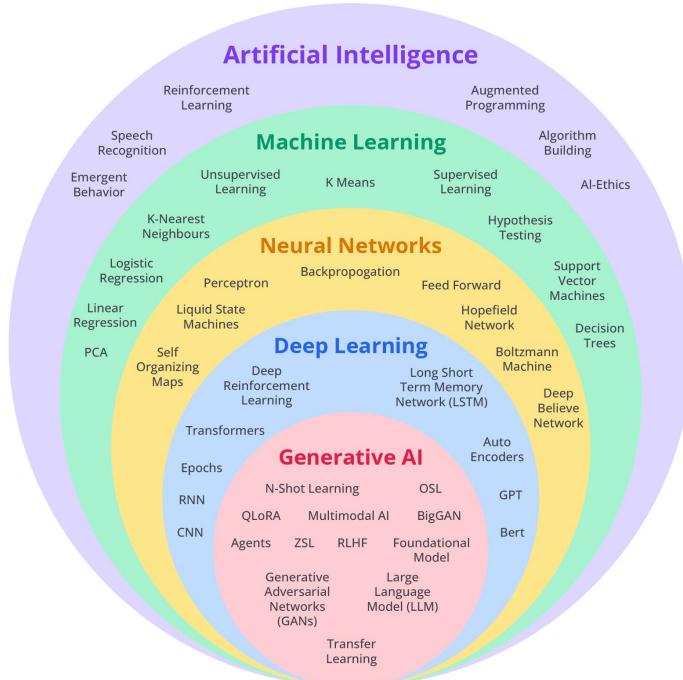
3



4

**Introduction to Time  
Series Prediction**

# Intermezzo



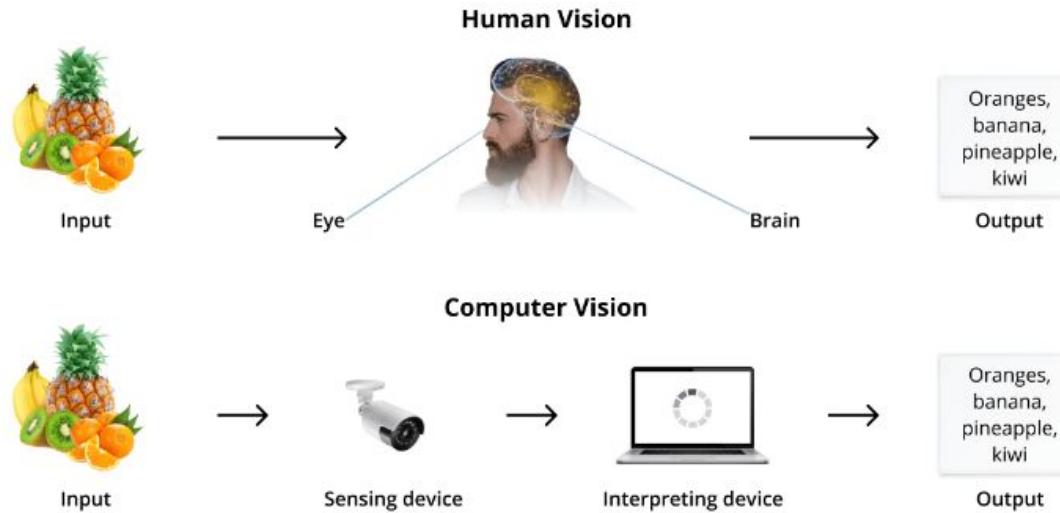
**AI:** Artificial Intelligence (AI) refers to the broad field of developing systems or machines that can mimic human intelligence to perform tasks such as decision-making, problem-solving, and understanding language.

**Machine Learning:** Machine Learning (ML) is a subset of AI where computers learn patterns from data to make predictions or decisions without being explicitly programmed.

**Deep Learning:** Deep Learning is a specialized area of ML that uses neural networks with many layers to analyze complex patterns in large datasets, such as images or speech.

**Generative AI:** Generative AI focuses on creating new content like text, images, or music by learning patterns from existing data and generating outputs that mimic or extend those patterns.

# Human Vision VS Computer Vision



# How Machine can Understand?



What humans see

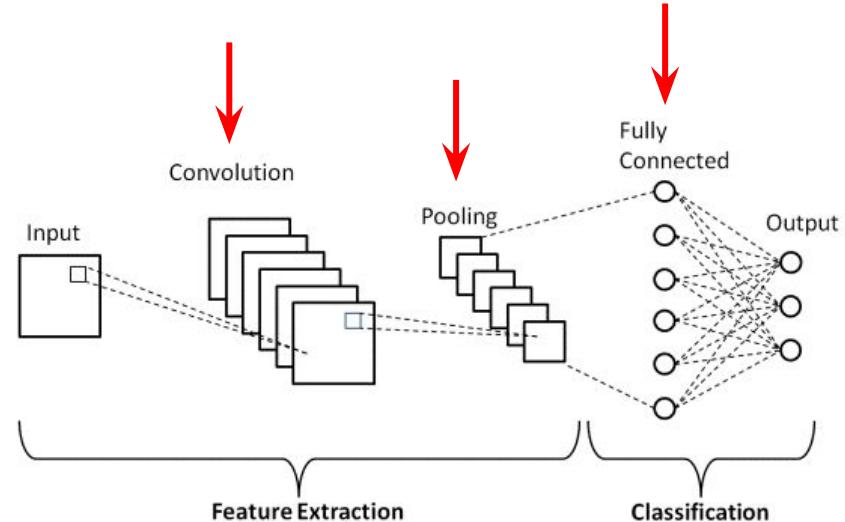
08 02 22 97 38 15 00 40 00 75 01 05 07 78 52 12 50 77 91 03  
19 49 99 40 17 81 18 57 60 67 17 40 98 43 69 45 04 56 62 00  
61 49 31 73 55 79 14 29 93 71 10 67 53 88 30 03 19 13 36 65  
52 70 96 23 04 60 11 42 69 24 65 66 01 32 56 71 37 02 36 92  
22 31 16 71 51 67 43 00 41 92 36 54 22 40 40 25 66 33 13 60  
24 47 32 60 99 03 45 02 44 75 33 55 70 36 04 20 35 17 12 50  
32 90 61 28 64 23 47 10 26 35 40 47 59 54 70 66 10 30 44 70  
67 26 20 68 02 62 12 20 95 63 94 39 63 08 40 91 66 49 94 22  
24 55 56 05 66 73 99 26 97 17 70 70 96 83 14 00 34 09 43 72  
21 36 23 09 75 00 76 44 20 45 35 14 00 61 33 97 34 31 33 95  
70 17 53 25 22 75 31 67 15 94 03 80 40 62 16 14 09 53 56 92  
16 39 06 42 96 35 31 47 55 53 88 24 00 17 54 24 36 29 85 57  
26 36 00 48 55 71 69 07 05 44 44 37 44 60 21 55 51 54 17 55  
19 60 01 65 05 94 47 69 20 73 92 13 86 52 17 77 04 89 55 40  
04 52 02 03 97 35 99 16 07 97 57 32 16 26 26 79 33 27 98 66  
00 36 63 07 57 62 20 72 03 46 33 67 46 55 22 32 63 93 53 69  
04 42 16 73 38 25 39 11 24 94 72 10 08 46 29 32 40 42 76 36  
20 49 36 41 72 30 23 00 34 62 99 69 02 67 59 05 74 04 36 16  
20 73 33 29 78 31 90 01 74 31 49 71 48 86 01 16 23 57 05 54  
01 70 54 71 03 51 54 69 16 92 33 48 61 43 52 01 89 19 67 48

What the computer sees

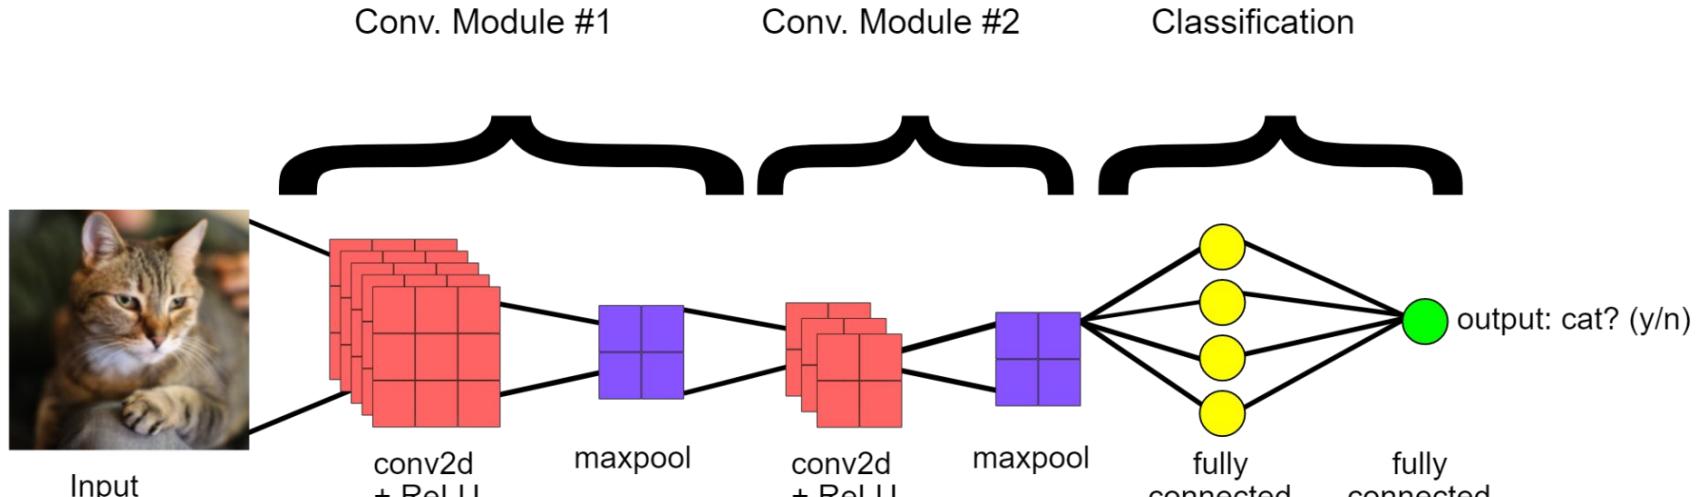
# **Image Classification with Convolutional Neural Network**

# What is Convolutional Neural Network?

CNN is a type of neural network model which allows us to extract higher representations for the image content.



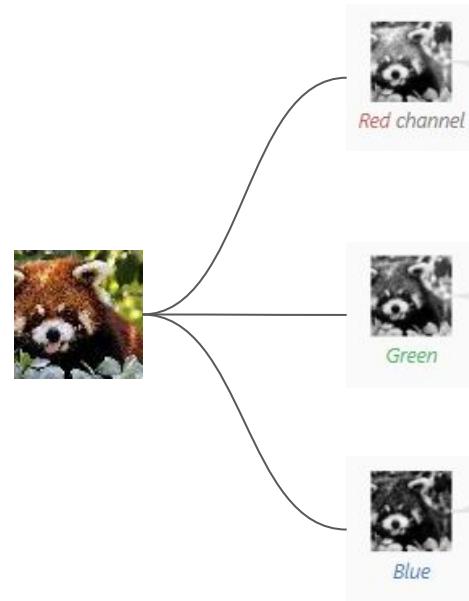
# Typical CNN Architecture



**Shape:** (32, 32, 3)

# Typical CNN Architecture - **Input Layer**

The input layer functions to receive input data or input from the dataset that will be processed by the network.



# Typical CNN Architecture - Convolution Layer

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

(3,3)  
Stride = 1

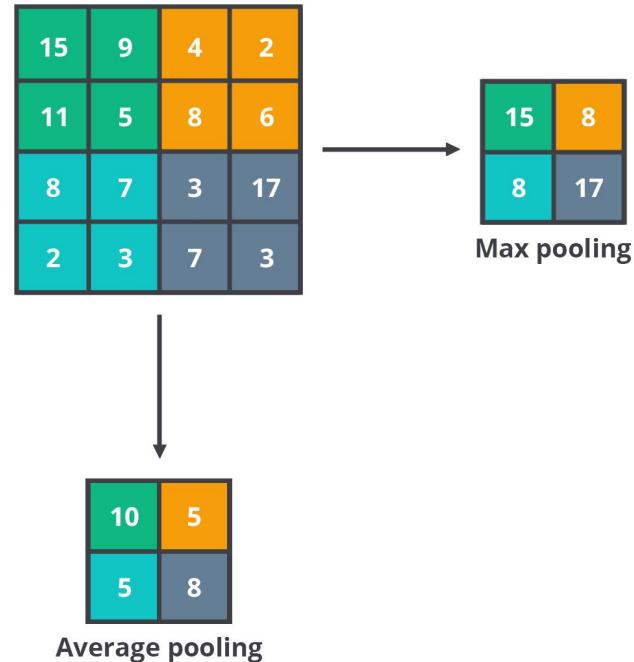
$$\begin{matrix} & \ast & \\ \begin{matrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{matrix} & & \begin{matrix} 6 \\ \vdots \\ \vdots \end{matrix} \end{matrix} =$$

$7x1 + 4x1 + 3x1 +$   
 $2x0 + 5x0 + 3x0 +$   
 $3x-1 + 3x-1 + 2x-1$

# Typical CNN Architecture - Pooling Layer

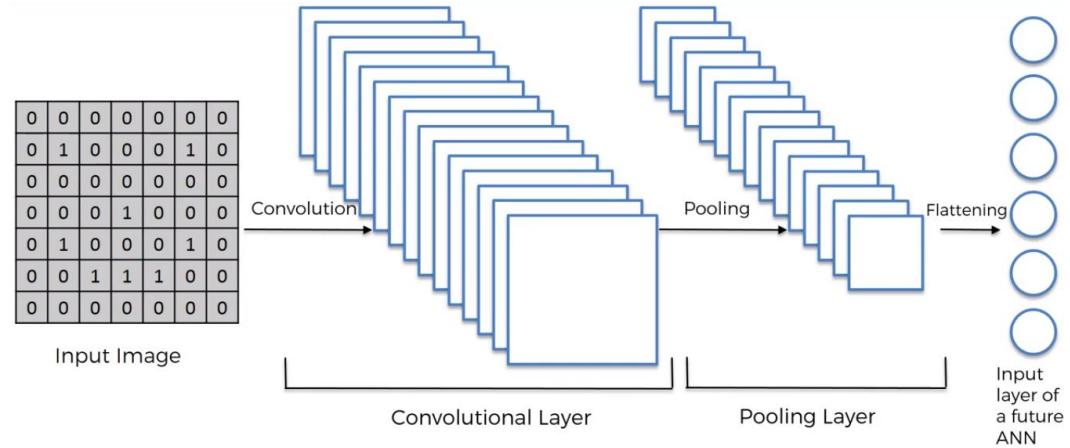
Pooling is a **type of downsampling** that often occurs after convolution.

The goal is to **reduce the size** of the training data before it goes into the fully connected network without losing much information.



# Typical CNN Architecture - Flatten Layer

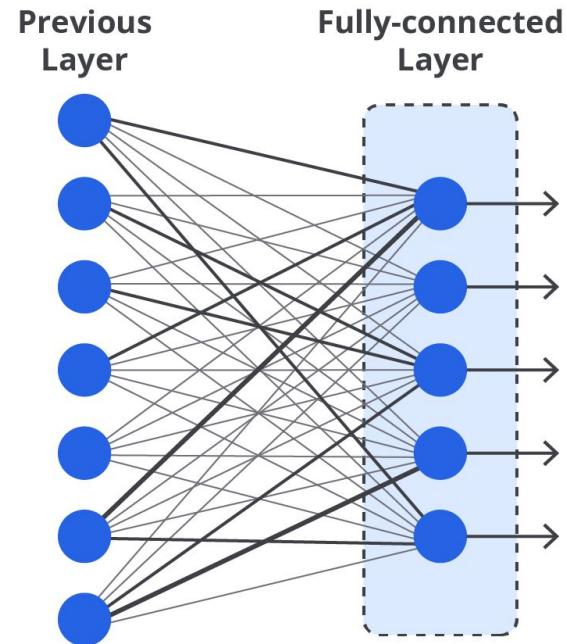
The Flatten layer **reduces the dimensionality** of the data, which can help reduce the number of parameters in the subsequent fully connected layers.



# Typical CNN Architecture - Fully Connected Layers

Fully connected layers are an important part of a neural network located at the end of the architecture, usually after a series of convolution layers and pooling layers in a CNN.

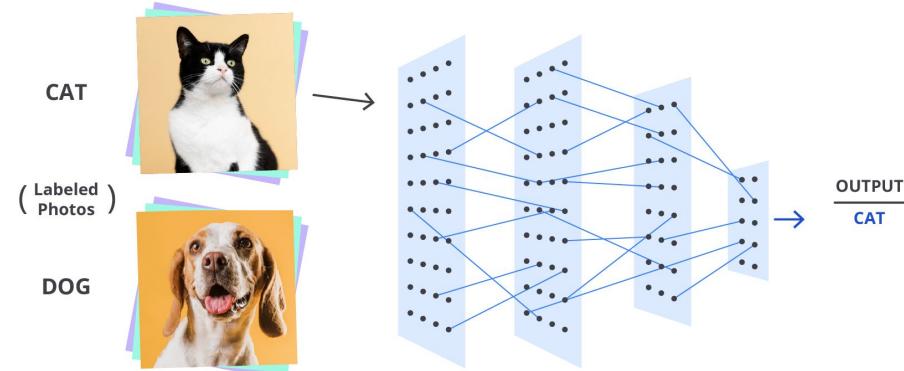
The function of this layer is to combine all the features extracted from the previous layers to make a final decision, such as classification or prediction.



# Typical CNN Architecture - Output Layers

Output layers, the last part, produce the predictions or final output of the model.

The function of this layer is highly dependent on the type of task being performed by the network.



# How to build CNN Architecture?

```
## START CODE HERE ##

# Retrieve dimensions from A_prev's shape (≈1 line)
(m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

# Retrieve dimensions from W's shape (≈1 line)
(f, f, n_C_prev, n_C) = W.shape

# Retrieve information from "hparameters" (≈2 lines)
stride = hparameters['stride']
pad = hparameters['pad']

# Compute the dimensions of the CONV output volume using the formula given above. Hint: use int() to floor. (≈2 Lines)
n_H = int((n_H_prev - f + 2 * pad) / stride) + 1
n_W = int((n_W_prev - f + 2 * pad) / stride) + 1

# Initialize the output volume Z with zeros. (≈1 line)
Z = np.zeros((m, n_H, n_W, n_C))

# Create A_prev_pad by padding A_prev
A_prev_pad = zero_pad(A_prev, pad)

for i in range(m):
    a_prev_pad = A_prev_pad[i]
    for h in range(n_H):
        for w in range(n_W):
            for c in range(n_C):

                # Loop over the batch of training examples
                # Select ith training example's padded activation
                # Loop over vertical axis of the output volume
                # Loop over horizontal axis of the output volume
                # loop over channels (= #filters) of the output volume

                # Find the corners of the current "slice" (≈4 Lines)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Use the corners to define the (3D) slice of a_prev_pad (See Hint above the cell). (≈1 line)
                a_slice_prev = a_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :]

                # Convolve the (3D) slice with the correct filter W and bias b, to get back one output neuron. (≈1 Line)
                Z[i, h, w, c] = conv_single_step(a_slice_prev, W[...,c], b[...,c])
```

Convolutional

```
# Retrieve dimensions from the input shape
(m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

# Retrieve hyperparameters from "hparameters"
f = hparameters['f']
stride = hparameters['stride']

# Define the dimensions of the output
n_H = int(1 + (n_H_prev - f) / stride)
n_W = int(1 + (n_W_prev - f) / stride)
n_C = n_C_prev

# Initialize output matrix A
A = np.zeros((m, n_H, n_W, n_C))

## START CODE HERE ##
for i in range(m):
    for h in range(n_H):
        for w in range(n_W):
            for c in range(n_C):

                # Loop over the training examples
                # Loop on the vertical axis of the output volume
                # Loop on the horizontal axis of the output volume
                # Loop over the channels of the output volume

                # Find the corners of the current "slice" (≈4 Lines)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Use the corners to define the current slice on the ith training example of A_prev, channel c. (≈1 line)
                a_prev_slice = A_prev[i, vert_start:vert_end, horiz_start:horiz_end, c]

                # Compute the pooling operation on the slice. Use an if statement to differentiate the modes. Use np
                if mode == "max":
                    A[i, h, w, c] = np.max(a_prev_slice)
                elif mode == "average":
                    A[i, h, w, c] = np.mean(a_prev_slice)

## END CODE HERE ##
```

Max Pooling

# How to build CNN **Architecture**?

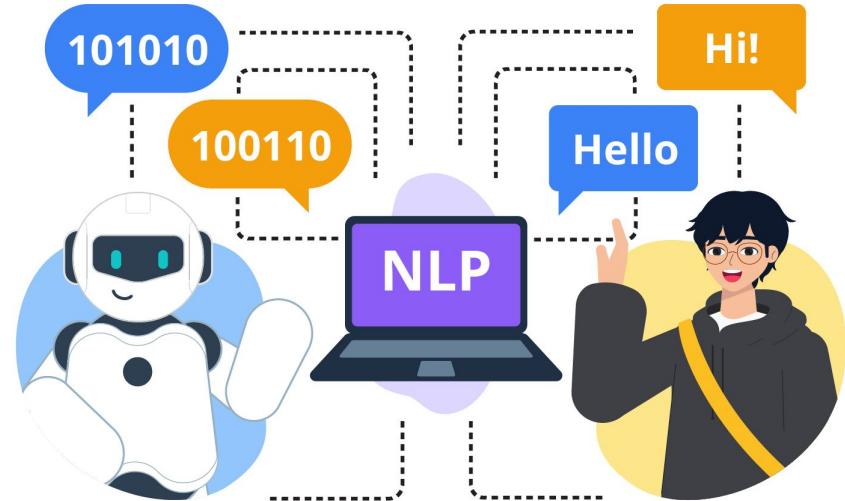


# [Hands-on] Unlocking the Power of CNNs for Image Classification

# An Overview of Natural Language Processing

## What is NLP?

Natural Language Processing (NLP) is a branch of Artificial Intelligence that gives machines the ability to **read**, **process** and **derive meaning** from human languages



# The Applications of NLP in Daily Life



Google Translate interface showing the translation of Indonesian text "Halo, saya sedang mempelajari pemrosesan bahasa alami. Salah satu kegunaannya adalah penerjemah otomatis seperti Google Terjemahan!" into Korean.

## Sentiment Analysis

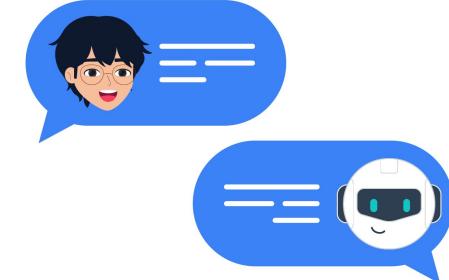


Menemukan opini, emosi, dan perasaan pengguna tentang suatu produk atau layanan

Our team has **less** project this quarter.

• GRAMMAR  
**less** → **fewer**

It appears that the quantifier **less** does not fit with the countable noun **projects**. Consider changing the quantifier or the noun.



# How Computer Processing A Text

I love my cat



{ "I" : 1, "love" : 2, "my" : 3, "cat" : 4 }



Encoded Word = [ 1, 2, 3, 4 ]



```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer

sentences = ["I love my cat"]

tokenizer = Tokenizer(num_words=100)
tokenizer.fit_on_texts(sentences)
sequences = tokenizer.texts_to_sequences(sentences)

print(tokenizer.word_index)
print(sequences)
```

# Additional Tips for How Computer Processing A Text

Case folding aims to **convert** all letters in a text document into **lowercase letters**.

Saya sedang belajar Natural Language Processing

# Additional Tips for How Computer Processing A Text

Removal special characters aims to **rid the text of unwanted special characters**, such as numbers, punctuation marks, symbols, or other special characters.

@dicoding: Hallo semuanya!! Masih semangat belajar NLP kan??! Berapa persen semangatnya? 50%? 60%? atau 100%?  
#menyalaabangkuh! >.<

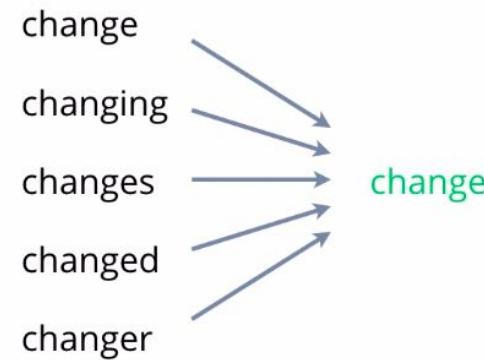
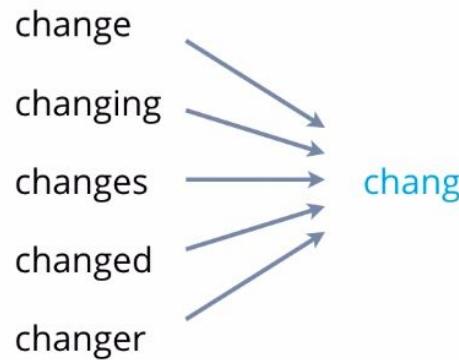
# Additional Tips for How Computer Processing A Text

Stopwords removal (filtering) is the **removal of stopwords or common words** that appear frequently, but do not have high informative value in text analysis.

Makanan favorit adik saya adalah nasi goreng dan ayam goreng crispy karena sangat enak dan menggugah selera

# Additional Tips for How Computer Processing A Text

## Stemming vs Lemmatization



# Additional Tips for How Computer Processing A Text

```
● ● ●  
import tensorflow as tf  
from tensorflow import keras  
from tensorflow.keras.preprocessing.text import Tokenizer  
  
sentences = [  
    "I love my cat",  
    "I love my dog",  
    "Do you think my dog is amazing?",  
    "Do you think my cat is cute?"  
]  
  
tokenizer = Tokenizer(num_words=100, oov_token="<00V>")  
tokenizer.fit_on_texts(sentences)  
sequences = tokenizer.texts_to_sequences(sentences)  
  
print(tokenizer.word_index)  
print(sequences)
```

```
[[3, 4, 2, 5], [3, 4, 2, 6], [7, 8, 9, 2, 6, 10, 11], [7, 8, 9, 2, 5, 10, 12]]
```

Put a special value in when an unseen word encountered.

# Padding

Padding is used to fit all dictionaries in one sequence. So, all sequences have the same length.

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

sentences = [
    "I love my cat",
    "I love my dog",
    "Do you think my dog is amazing?",
    "Do you think my cat is cute?"
]

tokenizer = Tokenizer(num_words=100, oov_token=<OOV>)
tokenizer.fit_on_texts(sentences)
sequences = tokenizer.texts_to_sequences(sentences)

padded = pad_sequences(sequences, padding='post', truncating='post', maxlen=10 )

print(tokenizer.word_index)
print(padded)
```

```
[[ 3 4 2 5 0 0 0 0 0 0]
 [ 3 4 2 6 0 0 0 0 0 0]
 [ 7 8 9 2 6 10 11 0 0 0]
 [ 7 8 9 2 5 10 12 0 0 0]]
```

# Word Embeddings

Word embeddings are a type of word representation that allows words with similar meanings to have **similar representations**.



# Word Embedding

An embedding is a dense vector of floating point values (the length of the vector is a parameter you specify).

## A 4-dimensional embedding

**cat** =>

1.2	-0.1	4.3	3.2
0.4	2.5	-0.9	0.5
2.1	0.3	0.1	0.4

**mat** =>

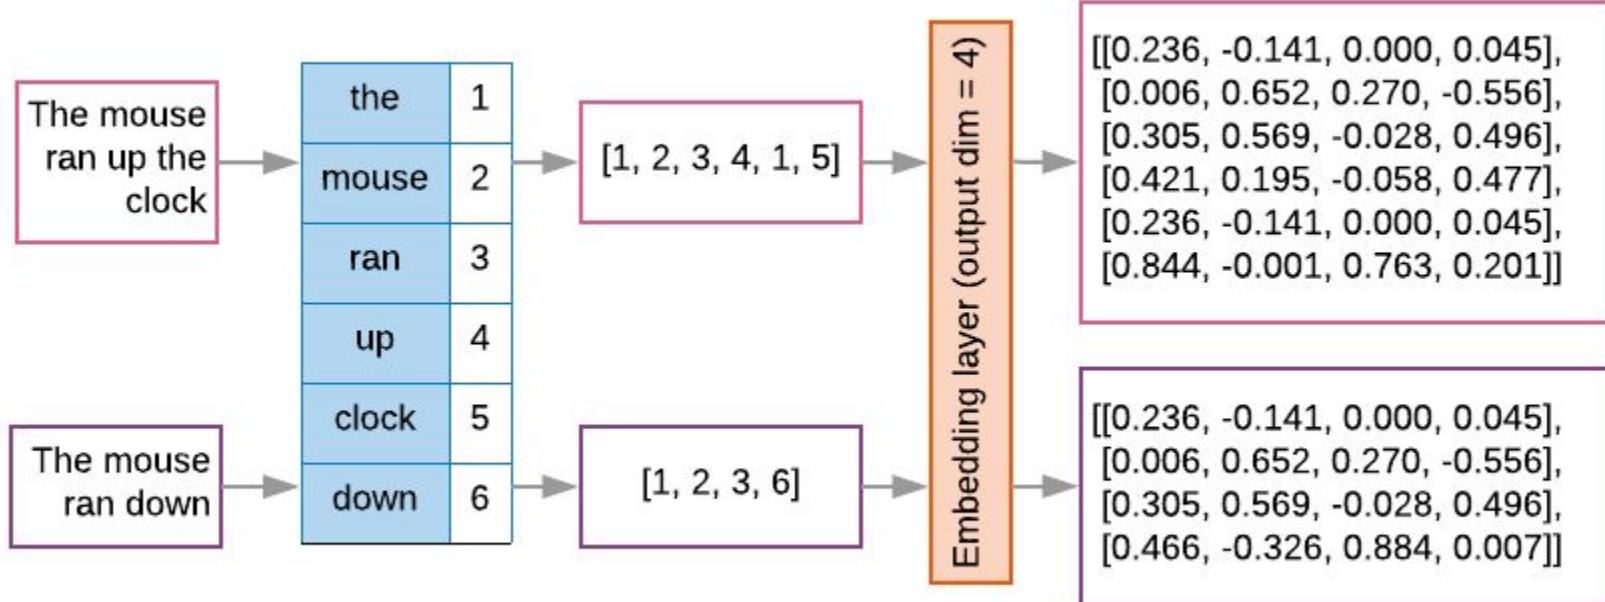
**on** =>

1.2	-0.1	4.3	3.2
0.4	2.5	-0.9	0.5
2.1	0.3	0.1	0.4

...

...

# Word Embedding



# [Hands-on] Riding the Wave of **NLP** Experience

# Break Time!

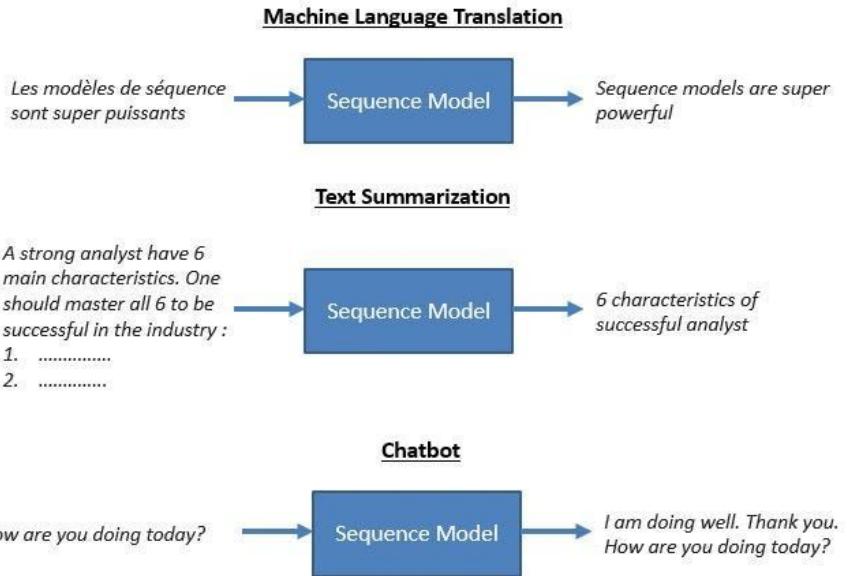


TERTURU

# **Exploring RNNs and LSTM**

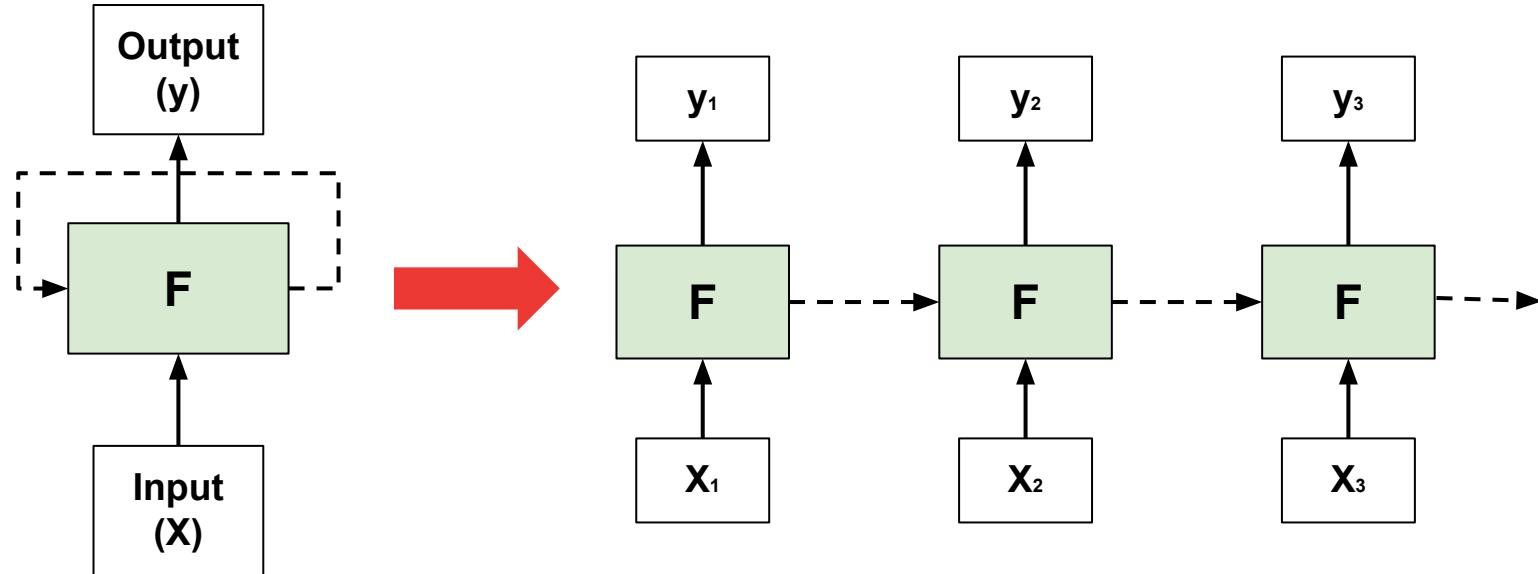
# Limitation of Basic ANN Model

- In sequential problems, the **order of values matters**.
- To handle this, the model needs to **use previous experience to process the next input value**.
- Basic ANN model **can't do this**.

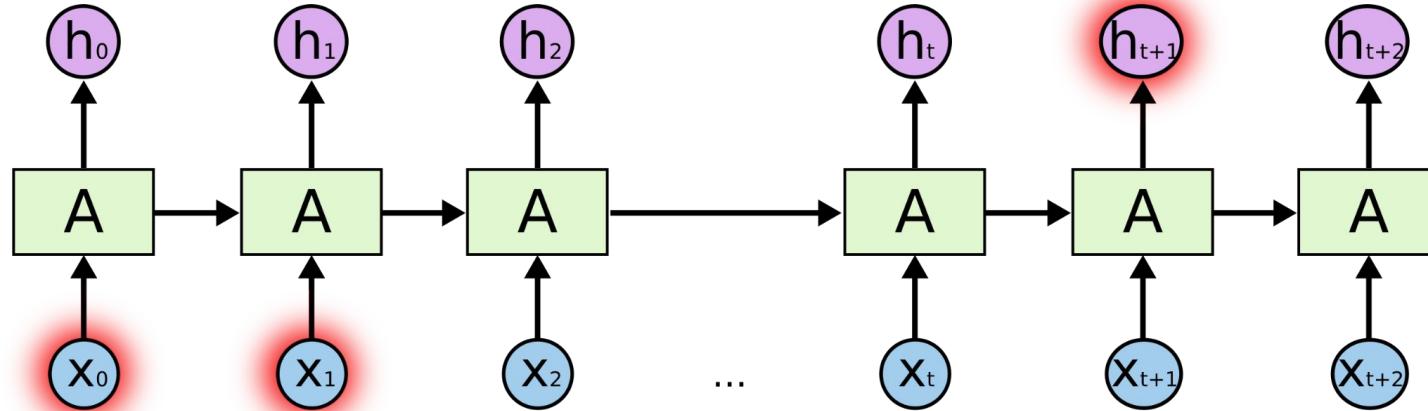


# RNN Model Architecture

A recurrent neural network (RNN) algorithm is a type of deep learning algorithm designed to process sequence data.



# The Downside of RNN Models



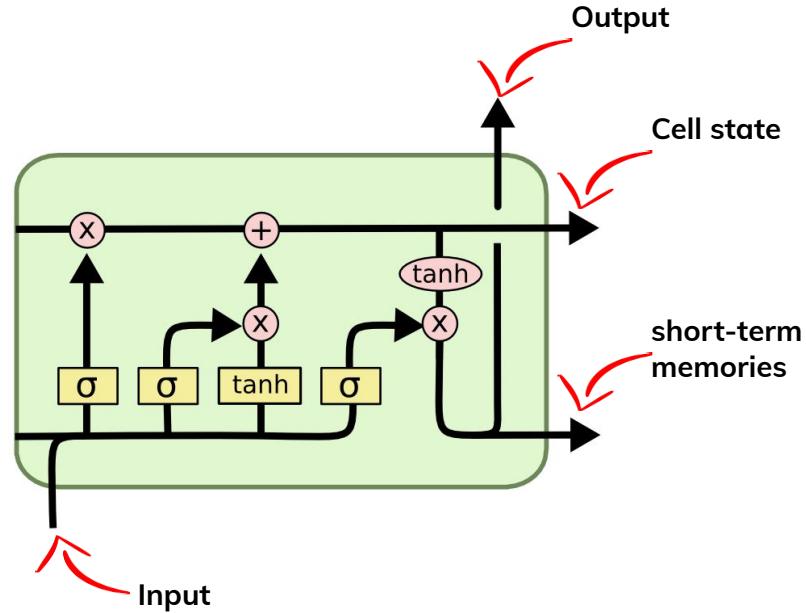
RNN never really did well at handling large sequences of data, like long paragraphs. By the time they were process the end of a paragraph they'd forget what happened in the beginning.

**Is it the only option to solve the problem?**

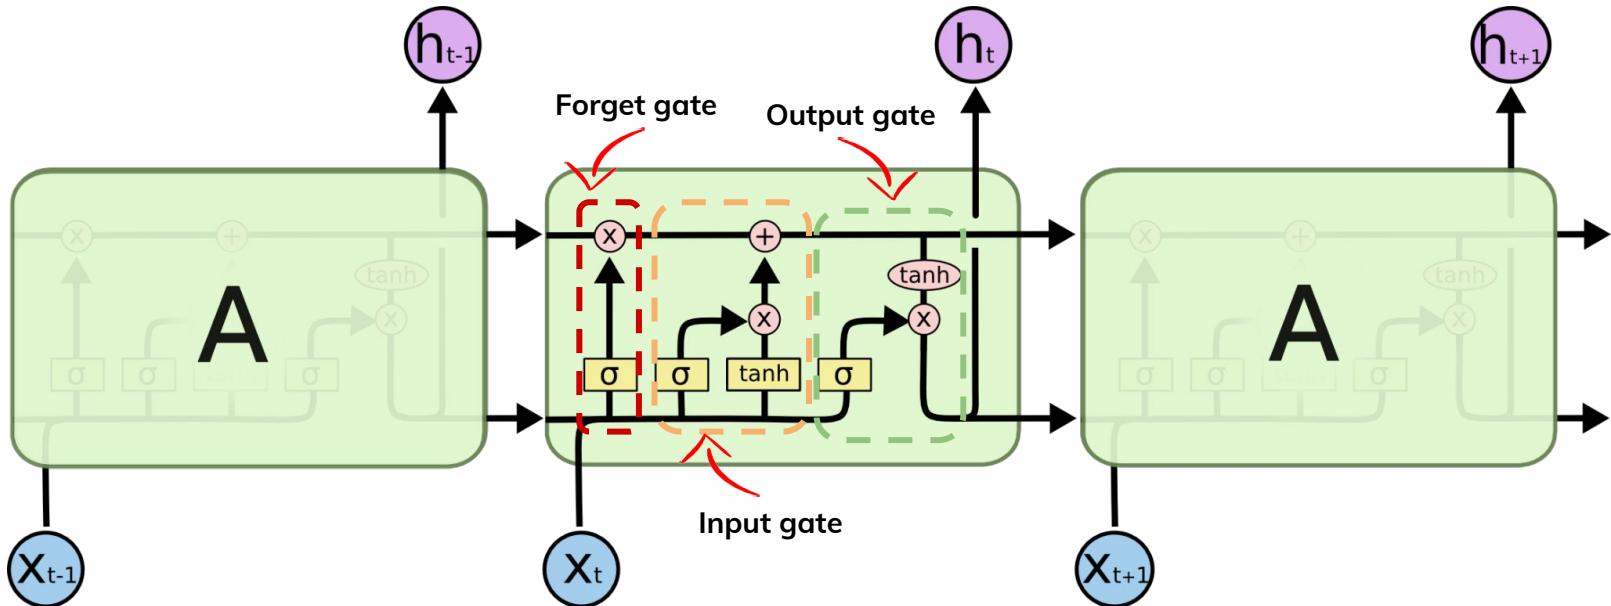


# Long Short Term Memory Network

- LSTMs are a **special kind of RNN**, capable of learning long-term dependencies.
- LSTM have **long-term memories called cell state** to handle long-term dependencies.



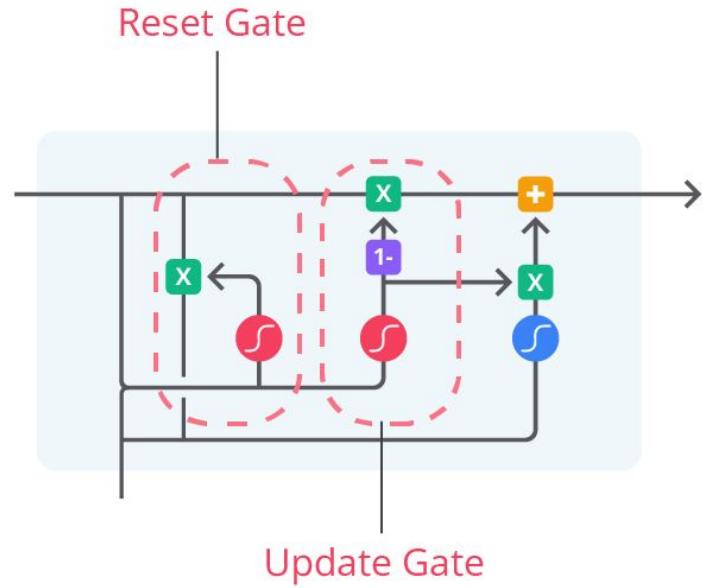
# LSTM Model Architecture



# Gated Recurrent Unit (GRU)

GRUs are a type of unit in RNNs developed to model long-range dependencies in **sequence data** in a **more efficient way** than previous models, such as LSTM units.  
GRUs have two main gate types.

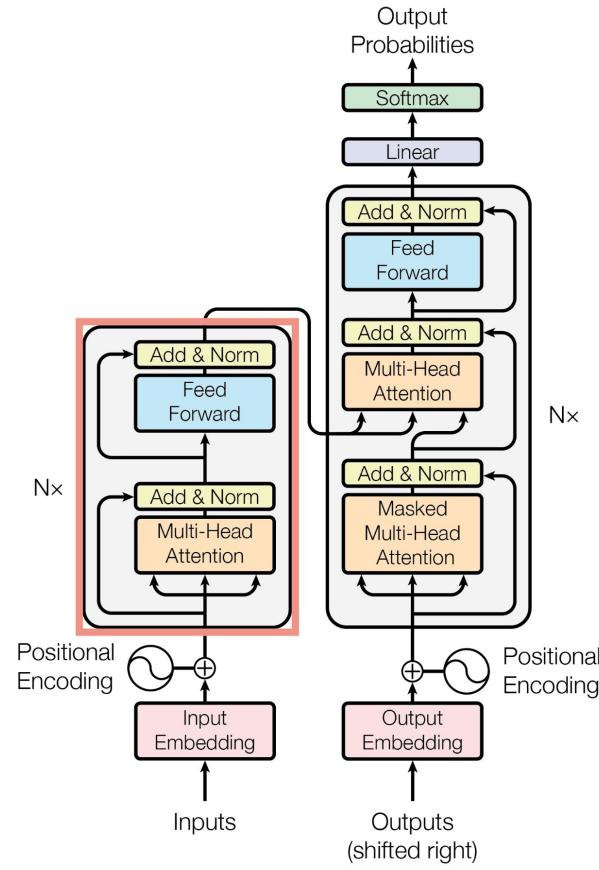
- Reset Gate
- Update Gate



# Sneak Peek Transformers

Transformers are a type of deep learning model used for natural language processing (NLP) and computer vision (CV) tasks. They utilize a mechanism called “**self-attention**” to process sequential input data.

Transformers can process the entire input data at once, capturing context and relevance.



# Sequence Prediction



Sequence prediction is a type of machine learning task that involves **predicting the next value or sequence** of values in a given series.

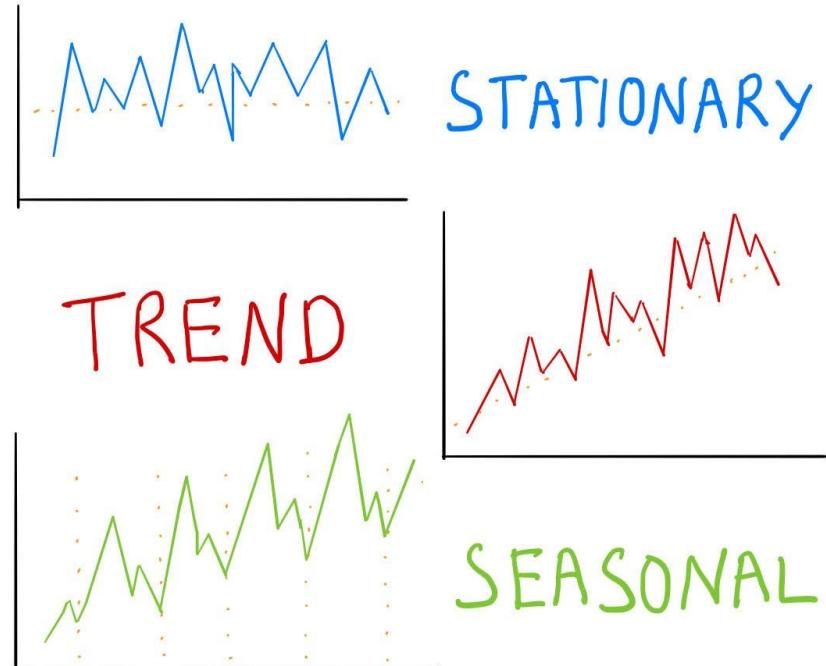
# Getting Started with Time Series

# What is Time Series?

An ordered sequence of values that are usually equally spaced over time.

What are they?

- Stock prices
- Weather forecast
- Historical trends



# Univariate & Multivariate Time Series

## Univariate

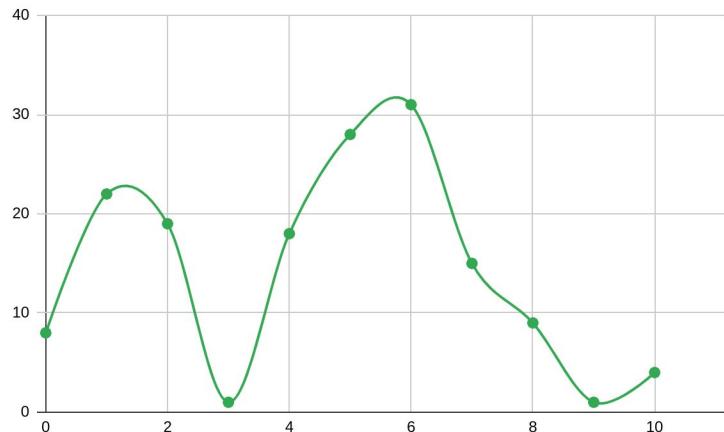
Time	Temperature
05.00 WIB	15 °C
06.00 WIB	15 °C
07.00 WIB	14 °C
08.00 WIB	14 °C
09.00 WIB	16 °C
10.00 WIB	17 °C
11.00 WIB	18 °C
12.00 WIB	19 °C
13.00 WIB	19 °C
14.00 WIB	21 °C
15.00 WIB	22 °C
16.00 WIB	22 °C
17.00 WIB	22 °C
18.00 WIB	21 °C
19.00 WIB	20 °C
20.00 WIB	18 °C
21.00 WIB	18 °C

## Multivariate

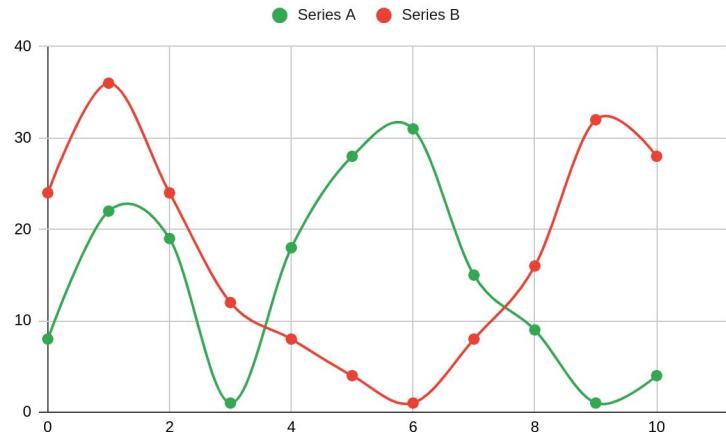
Time	Temperature	Cloud Cover	Dew Point	Humidity	Wind
5:00 am	15 °C	97%	10 °C	74%	8 mph SSE
6:00 am	15 °C	89%	10 °C	75%	8 mph SSE
7:00 am	14 °C	79%	10 °C	76%	7 mph SSE
8:00 am	14 °C	74%	10 °C	77%	7 mph S
9:00 am	16 °C	74%	10 °C	74%	7 mph S
10:00 am	17 °C	74%	11 °C	70%	8 mph S
11:00 am	18 °C	76%	11 °C	65%	8 mph SSW
12:00 am	19 °C	80%	11 °C	60%	8 mph SSW
1:00 pm	19 °C	78%	11 °C	58%	10 mph SW
2:00 pm	21 °C	71%	11 °C	54%	10 mph SW
3:00 pm	22 °C	75%	11 °C	52%	11 mph SW
4:00 pm	22 °C	78%	11 °C	52%	11 mph SW
5:00 pm	22 °C	78%	11 °C	52%	12 mph SW
6:00 pm	21 °C	78%	11 °C	54%	11 mph SW
7:00 pm	20 °C	87%	11 °C	60%	12 mph SW
8:00 pm	18 °C	100%	12 °C	66%	11 mph SSW
9:00 pm	18 °C	100%	12 °C	72%	13 mph SSW

# Univariate & Multivariate Time Series

## Univariate

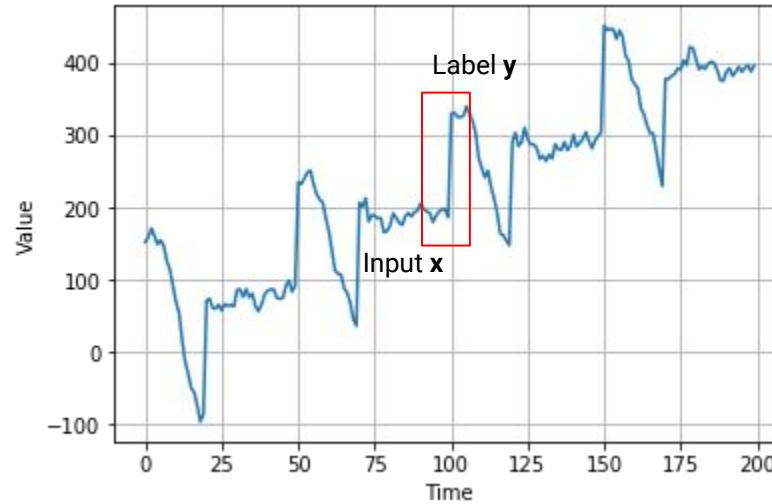


## Multivariate



# Preparing Features and Labels

Windowing the data to create the **features** and **labels**



# [Hands-on] Sailing through Time: A Navigator's Guide to Time Series

# Discussions



**NOOB BERTANYA,  
LORD MENJAWAB**

# Quiz and Feedback

BINTANG 5 PLIS



OH TUHAN



BERI HAMBA

# Conclusion

You have learned basic concepts in **NLP** and **time series prediction**. You also learn about advance deep learning model such as **CNN**, **RNN**, and **LSTM**.

These knowledge will give you a solid fundamental concept to becoming a **machine learning practitioner**.



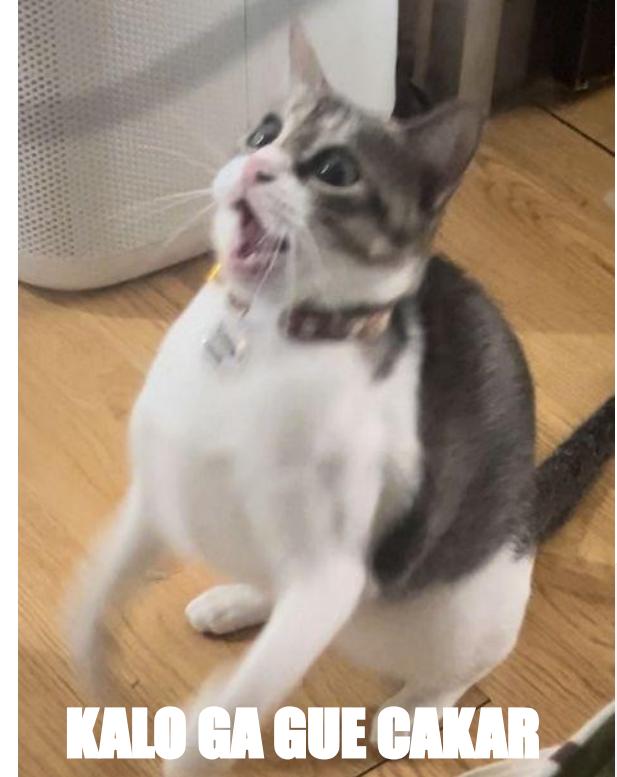
# Thank You



Follow majikan saia:  
Tiktok: @jenniethecatwoman

Currently at ~1500 Followers

 bangkit



KALO GA GUE CAKAR

LU